

Building and Running Apps on Kestrel

Timothy H. Kaiser, Ph.D.
tkaiser2@nrel.gov

ABSTRACT

This talk will cover the parallel and serial (MPI/C/Fortran) programming and execution environments on Kestrel. We'll approach this from a "toolchain" perspective. That is, we'll show makefiles which load the required modules then compile parallel (hybrid MPI/OpenMP) and serial Fortran and C codes. Then we'll show the commands to run the applications on compute nodes.

We'll tie this all together by presenting a script that will make and run example codes using several programming environments including PrgEnv-cray, PrgEnv-gnu, PrgEnv-intel, and IntelMPI, OpenMPI.

The final result of running this script will be output from programs for each programming environment.

We'll also discuss nuisances of the PrgEnv-* environments and some surprising interactions between Cray's gcc compiler modules and IntelMPI **and various Intel modules.**

One of the programs will be discussed in more detail and we'll show how it can be used to diagnose affinity issues.

ABSTRACT

The National Renewable Energy Lab has just launched (and submitted a top 500 Run) its newest HPC platform, Kestrel, with 2144 nodes with dual Intel Sapphire Rapids processors. We'll discuss MPI/OpenMP affinity mapping and testing. We'll present a hybrid MPI/OpenMP test code that reports affinity as a function of environmental settings, tasks, and threads. We show that without attentiveness to affinity performance can be adversely effected However, we'll show how to get ideal mapping, where tasks and threads are laid out for performance. We present a batch script that can be run to sweep over various command line, environmental settings and task/thread combinations. In addition to Intel compilers the script will test Cray, MPICH, and OpenMPI. We recommend the test code be run before a production run to ensure the desired mappings of tasks and threads. A git repository will be available with all codes and scripts.

CURRENT KESTREL CONFIGURATION

Number of Nodes	Processors	Memory	Accelerators	Local Storage
2144	Dual socket Intel Xeon Sapphire Rapids (52-core) processors	256 GB DDR5	N/A	256 nodes with 1.92 TB NVMe M.2
10	Dual socket Intel Xeon Sapphire Rapids (52-core) processors	2 TB DDR5	N/A	8 x 1.6 TB NVMe
8	Dual socket Intel Xeon Sapphire Rapids (52-core) processors	256 GB DDR5	2 NVIDIA A40 GPUs	2 x 3.84 TB NVMe

SCHEDULE

- Some links
- What's up with all the module paths
- Module tricks
- Compile/Run with Intel compilers and Intel MPI
- A bunch of Intel modules
- Prog-^{*}
 - What are they?
 - What is craype-x86-spr
 - More compiler wrappers
 - Compile/Run with Prog-^{*}
 - Interactions between Intel MPI and Prog-^{*}
- Affinity testing - What's it about and why is it important?
- Complete example script

SOME LINKS

- <https://nrel.github.io/HPC/Documentation/>
- <https://nrel.github.io/HPC/Documentation/Systems/Kestrel/Environments/>
- <https://nrel.github.io/HPC/Documentation/Systems/Kestrel/Environments/Toolchains/>

- **Examples:**

```
tar -xzf /nopt/nrel/apps/examples/recreate.tgz
```

OR

```
git clone git@github.nrel.gov:hpc-apps/kestrel-tds.git ; cd kestrel-tds/affinity/tutorial  
sbatch -A MYACCOUNT script
```

- **Near Future:** <https://github.com/NREL/HPC/tree/master/kestrel>
- <https://nrel.github.io/HPC/Documentation/Environment/shell/>

MODULE PATHS

```
module avail 2>&1 | grep "\-\" | sed s/\-//g | sort
/etc/modulefiles
/nopt/lmod/modulefiles/core
/nopt/lmod/modulefiles/mix_compilers
/nopt/nrel/apps/modules/default/application
/nopt/nrel/apps/modules/default/compilers_mpi
/nopt/nrel/apps/modules/default/utilities_libraries
/opt/cray/modulefiles
/opt/cray/pe/lmod/modulefiles/core
/opt/cray/pe/lmod/modulefiles/cpu/x86spr/1.0
/opt/cray/pe/lmod/modulefiles/craypetargets/default
/usr/share/lmod/lmod/modulefiles/Core
```

- Any path with nrel in the name was installed by us, contains programs, newer compilers, utilities...
- Everything else is part of the base system
- Any module with mixed in the name is "support" and should not be loaded directly
- Part of the complexity is related to interdependence of modules

Some module "tricks"

```
[tkaiser2@kl1 ~]$module list      # these are the default modules
```

Currently Loaded Modules:

```
1) craype-x86-spr      3) craype-network-ofi      5) cce/15.0.0      7) cray-dsmml/0.2.2      9) cray-libsci/22.12.1.1
2) libfabric/1.15.2.0  4) perftools-base/22.12.0  6) craype/2.7.19   8) cray-mpich/8.1.23    10) PrgEnv-cray/8.3.3
```

```
[tkaiser2@kl1 ~]$module purge
```

```
[tkaiser2@kl1 ~]$module list
```

No modules loaded

```
[tkaiser2@kl1 ~]$module load openmpi/4.1.5-gcc      gcc/13.1.0
```

```
[tkaiser2@kl1 ~]$module list
```

Currently Loaded Modules:

```
1) openmpi/4.1.5-gcc  2) gcc/13.1.0
```

```
[tkaiser2@kl1 ~]$module save myopen
```

Saved current collection of modules to: "myopen"

```
[tkaiser2@kl1 ~]$module purge
```

```
[tkaiser2@kl1 ~]$
```

```
[tkaiser2@kl1 ~]$module restore system      # this restores the default modules
```

Resetting modules to system default. Resetting \$MODULEPATH back to system default. All extra directories will be removed from \$MODULEPATH.

```
[tkaiser2@kl1 ~]$module list
```

Currently Loaded Modules:

```
1) craype-x86-spr      3) craype-network-ofi      5) cce/15.0.0      7) cray-dsmml/0.2.2      9) cray-libsci/22.12.1.1
2) libfabric/1.15.2.0  4) perftools-base/22.12.0  6) craype/2.7.19   8) cray-mpich/8.1.23    10) PrgEnv-cray/8.3.3
```

```
[tkaiser2@kl1 ~]$module restore myopen      # this restores the my set of modules
```

Restoring modules from user's myopen

```
[tkaiser2@kl1 ~]$module list
```

Currently Loaded Modules:

```
1) openmpi/4.1.5-gcc  2) gcc/13.1.0
```

```
[tkaiser2@kl1 ~]$
```


INTEL COMPILERS WITH INTELMPI

- Briefly discuss example programs
- Same compilers are available on Eagle/Swift/Vermilion and are "common"
- Go over a misconception about compiling with mpicc and mpif90
- Compilers generating warnings
- We'll
 - Show modules for builds and runs
 - Build commands
 - Run script extras
 - Example output
 - Makefile

OUR EXAMPLES

- **phostone.c** **fhostone.F90**

- Hello world on steroids
- Hybrid MPI / OpenMP
- Many command line options
- Will use options to print out a line for each MPI task and OpenMP thread along with the node and core on which it is running
- Will run for 7 seconds.

PHOSTONE.C OUTPUT

MPI VERSION Intel(R) MPI Library 2021.10 for Linux* OS

task	thread	node name	first task	# on node	core
0000	0000	X1005C4S5B0N0	0000	0000	0000
0000	0001	X1005C4S5B0N0	0000	0000	0001
0000	0002	X1005C4S5B0N0	0000	0000	0002
0000	0003	X1005C4S5B0N0	0000	0000	0003
0000	0004	X1005C4S5B0N0	0000	0000	0004
0000	0005	X1005C4S5B0N0	0000	0000	0005
0000	0006	X1005C4S5B0N0	0000	0000	0006
0000	0007	X1005C4S5B0N0	0000	0000	0007
0001	0000	X1005C4S5B0N0	0000	0001	0052
0001	0001	X1005C4S5B0N0	0000	0001	0053
0001	0002	X1005C4S5B0N0	0000	0001	0054
0001	0003	X1005C4S5B0N0	0000	0001	0055
0001	0004	X1005C4S5B0N0	0000	0001	0056
0001	0005	X1005C4S5B0N0	0000	0001	0057
0001	0006	X1005C4S5B0N0	0000	0001	0058
0001	0007	X1005C4S5B0N0	0000	0001	0059
0002	0000	X1005c4s6n0n0	0002	0000	0000
0002	0001	X1005c4s6n0n0	0002	0000	0001
0002	0002	X1005c4s6n0n0	0002	0000	0002
0002	0003	X1005c4s6n0n0	0002	0000	0003
0002	0004	X1005c4s6n0n0	0002	0000	0004
0002	0005	X1005c4s6n0n0	0002	0000	0005
0002	0006	X1005c4s6n0n0	0002	0000	0006
0002	0007	X1005c4s6n0n0	0002	0000	0007
0003	0000	X1005c4s6n0n0	0002	0001	0052
0003	0001	X1005c4s6n0n0	0002	0001	0053
0003	0002	X1005c4s6n0n0	0002	0001	0054
0003	0003	X1005c4s6n0n0	0002	0001	0055
0003	0004	X1005c4s6n0n0	0002	0001	0056
0003	0005	X1005c4s6n0n0	0002	0001	0057
0003	0006	X1005c4s6n0n0	0002	0001	0058
0003	0007	X1005c4s6n0n0	0002	0001	0059
total time		7.001			

- Intel MPI
- 2 MPI tasks / node
- 8 OpenMP threads
- Sorted and enhanced

OUR EXAMPLES

- **ppong.c**

- Measures bandwidth between MPI tasks as a function of message size
- Measures the MPI_Barrier rate
- Typically run with N tasks per node where N is the the number of cores
 - By default will run between every set of 2 MPI tasks $104 \times 103 = 10,712$ sets
 - The file "todo" gives a subset of tasks to test (2 tasks on first node and 2 on the second)

MPI VERSION Intel(R) MPI Library 2019 Update 10 for Linux* OS

calling MPI_Send - MPI_Recv

1 x1000c1s0b0n0 4.761803e-10 1.392327e-07

2 x1000c1s0b0n1 4.761803e-10 1.725275e-07

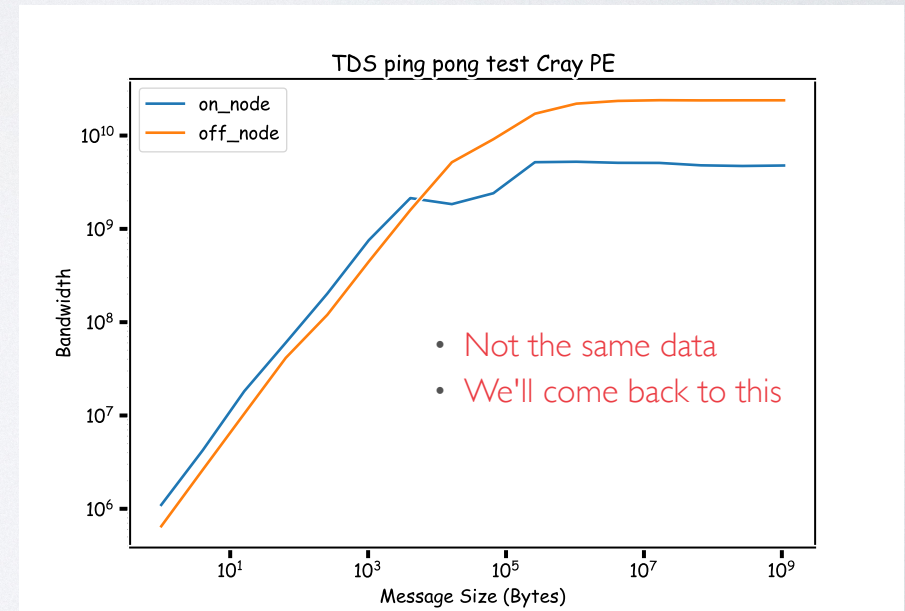
3 x1000c1s0b0n1 4.761802e-10 9.615906e-08

S	R	Size	Min Time	Ave Time	Max Time	Bandwidth	#
0	1	1	7.519964e-07	1.267607e-06	7.029879e-06	2.6596e+06	200
0	1	4	7.567462e-07	9.207008e-07	1.448463e-06	1.0572e+07	200
0	1	16	7.586554e-07	9.269662e-07	1.387508e-06	4.2180e+07	200
0	1	64	7.942552e-07	1.001559e-06	2.061017e-06	1.6116e+08	200
0	1	256	8.448493e-07	1.288456e-06	1.935754e-06	6.0603e+08	200
0	1	1024	1.669955e-06	1.944235e-06	3.213948e-06	1.2264e+09	200
0	1	4096	3.240118e-06	3.390790e-06	4.044105e-06	2.5283e+09	200
0	1	16384	8.213287e-06	8.448832e-06	9.251083e-06	3.9896e+09	200
0	1	65536	3.007229e-05	3.047969e-05	3.297967e-05	4.3586e+09	200
0	1	262144	5.973231e-05	6.045180e-05	7.064303e-05	8.7773e+09	200
0	1	1048576	2.803241e-04	2.858262e-04	3.947082e-04	7.4812e+09	200
0	1	4194304	1.366746e-03	1.457995e-03	1.646080e-03	6.1376e+09	200
0	1	16777216	5.132919e-03	5.218553e-03	5.361979e-03	6.5371e+09	97
0	1	67108864	2.027069e-02	2.048386e-02	2.063082e-02	6.6213e+09	26
0	1	268435456	8.050813e-02	8.093094e-02	8.161121e-02	6.6685e+09	8
0	1	1073741824	3.237138e-01	3.242401e-01	3.245610e-01	6.6339e+09	3
...							
0	3	1	2.814471e-05	5.013239e-05	4.272152e-03	7.1061e+04	200
0	3	4	2.833414e-05	2.902550e-05	3.022437e-05	2.8234e+05	200
0	3	16	2.806301e-05	2.900259e-05	3.033981e-05	1.1403e+06	200
0	3	64	2.811824e-05	2.896586e-05	3.021923e-05	4.5522e+06	200
0	3	256	3.016659e-05	3.136935e-05	3.362375e-05	1.6972e+07	200
0	3	1024	3.418040e-05	3.728393e-05	4.138465e-05	5.9917e+07	200
0	3	4096	3.088359e-05	3.221709e-05	3.544362e-05	2.6525e+08	200
0	3	16384	3.776823e-05	3.949602e-05	4.517029e-05	8.6761e+08	200
0	3	65536	8.354518e-05	8.565556e-05	8.813986e-05	1.5689e+09	200
0	3	262144	1.879028e-04	1.898664e-04	1.989481e-04	2.7902e+09	200
0	3	1048576	6.035536e-04	6.115581e-04	6.173814e-04	3.4747e+09	200
0	3	4194304	2.284192e-03	2.295765e-03	2.308413e-03	3.6725e+09	200
0	3	16777216	8.828382e-03	8.882622e-03	9.165776e-03	3.8007e+09	58
0	3	67108864	3.623108e-02	3.648658e-02	3.770787e-02	3.7045e+09	15
0	3	268435456	1.448475e-01	1.449654e-01	1.450861e-01	3.7065e+09	5
0	3	1073741824	5.784841e-01	5.787637e-01	5.790432e-01	3.7123e+09	2

...
Barriers/Second 45165.9
...

PPONG.C OUTPUT

- Intel MPI
- 2 MPI tasks / node
- Only showing info for two task pairs (0-1 & 0-3)



INTEL COMPILERS WITH INTELMPI

- Go over a misconception about compiling with mpicc and mpif90
- Compilers generating warnings

mpicc mpiicc mpif90 mpiifort

- mpicc - compile with gcc and Intel MPI
- mpiicc - compile with Intel icc and Intel MPI
- mpif90 - compile with gfortran and Intel MPI
- mpiifort - compile with Intel fortran and Intel MPI
- You can "force" compiles by other backends by setting `I_MPI_{CC,CXX,FC,F77,F90}`

WHAT ARE: ICX AND IFX?

You may see a warning when compiling with Intel compilers

`The Intel(R) Compiler Classic Compiler is deprecated...`

The Intel® oneAPI product packages provide two Fortran compilers. Intel Fortran Compiler Classic (ifort) provides best-in-class Fortran language features and performance for CPU. The Intel Fortran Compiler (ifx) enables developers needing OpenMP* offload to Intel GPUs. The OpenMP* 5.0, 5.1 GPU offload features in ifx are not available in ifort. For now ifort continues to be our best-in-class Fortran compiler for customers not needing GPU offload support.

- Notes:
 - Same for icc and icx except icc is being **replaced** by icx
 - Offload only works directly for Intel GPUs not Nvidia
(Should be a way to get this two work at least for C++)
 - Warning can be suppressed by export `-diag-disable=1044`

MODULE LOADS

For Builds

```
module load intel-oneapi-compilers  
module load intel-oneapi-mpi  
module load gcc/13.1.0
```

- The module load gcc is optional
- Gives you an newer version of gcc
- Can also load gcc/10.1.0
- Don't load other versions of gcc after loading the Intel modules (More on this later.)

For Running

```
module purge  
module load libfabric
```

- Usually don't need to load Intel modules
- libfabric gives access to the network

OUR BUILDS

1. Fortran with: Intel MPI and Intel Fortran compiler
2. C with: Intel MPI and Intel C compiler, older compiler (icc)
3. C with: Intel MPI and Intel C compiler, newer compiler (icx)
4. Fortran with: Intel MPI with gfortran Fortran compiler
5. C with: Intel MPI with gcc C compiler

OUR BUILDS

1. Fortran with: Intel MPI and Intel Fortran compiler

```
mpiifort -O3 -g -fopenmp ex1.f90
```

2. C with: Intel MPI and Intel C compiler, older compiler (icc)

```
mpiicc -O3 -g -fopenmp ex1.c
```

3. C with: Intel MPI and Intel C compiler, newer compiler (icx)

```
export I_MPI_CC=icx
```

```
mpiicc -O3 -g -fopenmp ex1.c
```

4. Fortran with: Intel MPI with gfortran Fortran compiler

```
mpif90 -O3 -g -fopenmp ex1.f90
```

5. C with: Intel MPI with gcc C compiler

```
mpicc -O3 -g -fopenmp ex1.c
```


MAKE

- Our example at <https://github.com/NREL/HPC/blob/master/kestrel/Toolchains/Code/Makefiles/Intel/makefile> builds for both Intel and gnu backends
- Here we just do Intel backends
- Or makefile uses a trick
 - Its default target is "recurse"
 - Making "recurse" loads module then calls make again for the actual targets

makefile

Intel MPI

```
SHELL:=/usr/bin/bash

recurse:
    module purge                ; \
    module load intel-oneapi    ; \
    module load intel-oneapi-mpi ; \
    module load gcc/13.1.0      ; \
    $(MAKE) -f $(firstword $(MAKEFILE_LIST)) both

both: f.impi c.impi pp.impi

#defines USEFAST
include makefile.include

ifeq ($(USEFAST),yes)
OPS=-DUSEFAST
EXTRA=getcore.o
endif

F90=mpiifort
CC=mpiicc

f.impi: fhostone.F90 $(EXTRA)
    $(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.impi
    rm -f getcore.o

c.impi: phostone.c
    $(CC) $(OPS) -fopenmp phostone.c -O3 -o c.impi

pp.impi: ppong.c
    $(CC) $(OPS) $(WES) ppong.c -O3 -o pp.impi

clean:
    rm -rf *o *mod* f.impi c.impi pp.impi
```


RUN COMMANDS

- Since phostone and fhostone are hybrid MPI/OpenMP codes we will set OpenMP variables.
 - # threads
 - # thread binding
- We'll also add some options to the srun line ensure we are getting good mapping of tasks and threads to cores
- We can run with various numbers of MPI tasks per node and number of OpenMP threads per task but the two multiplied together should not exceed then number of cores on a node (104).

Our script sets these openmp related variables. The first is familiar. KMP_AFFINITY is unique to Intel compilers. In this case we are telling the OS to scatter (spread) out our threads. OMP_PROC_BIND=spread does the same thing but it is not unique to Intel compilers. So in this case KMP_AFFINITY is actually redundant.

```
export OMP_NUM_THREADS=3  
export KMP_AFFINITY=scatter  
export OMP_PROC_BIND=spread
```

The next line

```
export BIND="--cpu-bind=v,cores"
```

is not technically used as an environmental variable but it will be used to create the srun command line (In later versions of the script). Passing --cpu-bind=v to srun will casue it to report threading information. The "cores" option tells srun to "Automatically generate masks binding tasks to cores." There are many other binding options as described in the srun man page. This setting works well for many programs.

Our srun command line options for 2 tasks per node and 3 threads per task are:

```
--mpi=pmi2 --cpu-bind=v,cores --threads-per-core=1 --tasks-per-node=2 --cpus-per-task=3
```

- --mpi=pmi2 : tells srun to use a particular launcher
- --cpu-bind=v,cores : discussed above
- --threads-per-core=1 : don't allow multiple threads to run on the same core. Without this option it is possible for multiple threads to end up on the same core, decreasing performance.
- --cpus-per-task=3 : The cpus-per-task should always be equal to OMP_NUM_THREADS.

SIMPLE SBATCH SCRIPT

```
#!/usr/bin/bash
#SBATCH --job-name="impi"
#SBATCH --nodes=2
#SBATCH --exclusive
#SBATCH --export=ALL
#SBATCH --time=00:10:00
#SBATCH --partition=standard

make -f makeimpi

export OMP_NUM_THREADS=3
export OMP_PROC_BIND=spread

CLA="-i -F -E -t 7"
# -i    : Print MPI_Init times for each task at end of run.
# -F    : Add columns to tell first MPI task on a node and and
#        the numbering of tasks on a node.
# -E    : Print thread info at 'E'nd of the run
# -t 7  : Run for 7 seconds

srun --mpi=pmi2 --cpu-bind=v,cores --threads-per-core=1 --tasks-per-node=2 --cpus-per-task=$OMP_NUM_THREADS ./f.impi $CLA > f.out
srun --mpi=pmi2 --cpu-bind=v,cores --threads-per-core=1 --tasks-per-node=2 --cpus-per-task=$OMP_NUM_THREADS ./c.impi $CLA > c.out
```


FORTTRAN - OUTPUT

```
[tkaiser2@kl1 src]$cat f.out
MPI Version: Intel(R) MPI Library 2021.10 for Linux* OS
```

task	thread	node name	first task	# on node	core
total time		7.00			
0000	0000	X1006C0S0B1N1	0000	0000	000
0000	0001	X1006C0S0B1N1	0000	0000	001
0000	0002	X1006C0S0B1N1	0000	0000	002
0001	0000	X1006C0S0B1N1	0000	0001	052
0001	0001	X1006C0S0B1N1	0000	0001	053
0001	0002	X1006C0S0B1N1	0000	0001	054
0002	0000	X1006C0S1B0N0	0002	0000	000
0002	0002	X1006C0S1B0N0	0002	0000	002
0002	0001	X1006C0S1B0N0	0002	0000	001
0003	0000	X1006C0S1B0N0	0002	0001	052
0003	0001	X1006C0S1B0N0	0002	0001	053
0003	0002	X1006C0S1B0N0	0002	0001	054
mpi_init	0	1696723552.3310	1696723552.9667	0.6357	
mpi_init	1	1696723552.3310	1696723552.9666	0.6356	
mpi_init	2	1696723552.3838	1696723552.9746	0.5908	
mpi_init	3	1696723552.3838	1696723552.9746	0.5908	

```
[tkaiser2@kl1 src]$
```


C - OUTPUT

```
[tkaiser2@kl1 src]$cat c.out
MPI VERSION Intel(R) MPI Library 2021.10 for Linux* OS
```

task	thread	node name	first task	# on node	core
total time 7.010					
0000	0000	X1006C0S0B1N1	0000	0000	0000
0000	0002	X1006C0S0B1N1	0000	0000	0002
0000	0001	X1006C0S0B1N1	0000	0000	0001
0001	0000	X1006C0S0B1N1	0000	0001	0052
0001	0001	X1006C0S0B1N1	0000	0001	0053
0001	0002	X1006C0S0B1N1	0000	0001	0054
0002	0000	X1006C0S1B0N0	0002	0000	0000
0002	0001	X1006C0S1B0N0	0002	0000	0001
0002	0002	X1006C0S1B0N0	0002	0000	0002
0003	0001	X1006C0S1B0N0	0002	0001	0053
0003	0002	X1006C0S1B0N0	0002	0001	0054
0003	0000	X1006C0S1B0N0	0002	0001	0052
mpi_init 0	1696723560.2818	1696723560.8412		0.5594	
mpi_init 1	1696723560.2824	1696723560.8408		0.5584	
mpi_init 2	1696723560.2902	1696723560.8409		0.5507	
mpi_init 3	1696723560.2902	1696723560.8409		0.5508	

```
[tkaiser2@kl1 src]$
```


PrgEnv-*

PrgEnv-amd/8.3.3

PrgEnv-aocc/8.3.3

PrgEnv-cray-amd/8.3.3

PrgEnv-cray/8.3.3

PrgEnv-gnu-amd/8.3.3

PrgEnv-gnu/8.3.3

PrgEnv-intel/8.3.3

PrgEnv-nvhpc/8.3.3

PrgEnv-nvidia/8.3.3

- Modules for using Cray's MPI with various backend compilers
- Red one currently work (cray, gnu, intel)
- Blue ones are for AMD processors and Nvidia GPUs (coming)
- PrgEnv-cray/8.3.3 is the default
- **All of these use the same MPI Library (Cray-MPICH)**

PrgEnv-cray loads several other modules

```
[tkaiser2@kl1 src]$module purge  
[tkaiser2@kl1 src]$module load PrgEnv-cray/8.3.3  
[tkaiser2@kl1 src]$module load craype-x86-spr  
[tkaiser2@kl1 src]$module list
```

Currently Loaded Modules:

1) cce/15.0.0	3) cray-dsmml/0.2.2	5) craype-network-ofi	7) cray-libsci/22.12.1.1	9) craype-x86-spr
2) craype/2.7.19	4) libfabric/1.15.2.0	6) cray-mpich/8.1.23	8) PrgEnv-cray/8.3.3	

These are the default modules, including crape-x86-spr

What is craype-x86-spr?

- x86-spr stands for Intel Sapphire Rapids - processors we have on Kestrel
- Loading this module sets environmental variables to allow certain optimizations for the Sapphire Rapids processors
- There are other choices that will not be useful until we get our GPU nodes

NEW COMPILER WRAPPERS

- Most MPI compilers mpicc, mpif90... are more or less wrapper scripts that call the underlying C or Fortran compilers with settings to point to the libraries
- ProgEnv-* takes this on step further
 - ftn = call underlying Fortran compiler and "auto detect" if it is a MPI program and build as such
 - cc = call underlying C compiler and "auto detect" if it is a MPI program and build as such
 - CC = call underlying C++ compiler and "auto detect" if it is a MPI program and build as such

PrgEnv-* wrappers

PrgEnv-*	Fortran (ftn)	C (cc)	MPI
PrgEnv-cray	Cray fortran (ftn)	cc - Clang based	Cray MPICH
PrgEnv-gnu	gfortran	gcc	Cray MPICH
PrgEnv-intel	ifort	icc	Cray MPICH

makefile

PrgEnv-cray

```
SHELL:=/usr/bin/bash

recurse:
    module purge                ; \
    module load craype-x86-spr  ; \
    module load PrgEnv-cray     ; \
    $(MAKE) -f $(firstword $(MAKEFILE_LIST)) both

both: f.cray c.cray pp.cray

#defines USEFAST
include makefile.include

ifeq ($(USEFAST),yes)
OPS=-DUSEFAST
EXTRA=getcore.o
endif

F90=ftn
CC=cc

f.cray: fhostone.F90 $(EXTRA)
    $(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.cray
    rm -f getcore.o

c.cray: phostone.c
    $(CC) $(OPS) -fopenmp phostone.c -O3 -o c.cray

pp.cray: ppong.c
    $(CC) $(OPS) $(WES) ppong.c -O3 -o pp.cray

clean:
    rm -rf *o *mod* f.cray c.cray pp.cray
```


Others are the nearly the same

```
tkaiser2-37907s:affinity tkaiser2$ diff makeprgcray makeprgintel
```

```
6c6,8
< module load PrgEnv-cray          ; \
---
> module load intel                ; \
> module load PrgEnv-intel         ; \
9c11
< both: f.cray c.cray pp.cray
---
> both: f.intel c.intel pp.intel
21a223,25
> f.intel: fhostone.F90 $(EXTRA)
> $(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.intel
> rm -f getcore.o
23,25c27,28
< f.cray: fhostone.F90 $(EXTRA)
< $(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.cray
< rm -f getcore.o
---
> c.intel: phostone.c
> $(CC) $(OPS) -fopenmp phostone.c -O3 -o c.intel
27,28c30,31
< c.cray: phostone.c
< $(CC) $(OPS) -fopenmp phostone.c -O3 -o c.cray
---
> pp.intel: ppong.c
> $(CC) $(OPS) $(WES) ppong.c -O3 -o pp.intel
30,33d32
< pp.cray: ppong.c
< $(CC) $(OPS) $(WES) ppong.c -O3 -o pp.cray
<
<
35c34
< rm -rf *o *mod* f.cray c.cray pp.cray
---
> rm -rf *o *mod* f.intel c.intel pp.intel
```

- We load different modules
- Compiler dependent options can change

makefile

PrgEnv-gnu

```
SHELL:=/usr/bin/bash

recurse:
    module purge                ; \
    module load craype-x86-spr  ; \
    module load PrgEnv-gnu      ; \
    module load gcc
    $(MAKE) -f $(firstword $(MAKEFILE_LIST)) both

both: f.gnu c.gnu pp.gnu

#defines USEFAST
include makefile.include

ifeq ($(USEFAST),yes)
OPS=-DUSEFAST
EXTRA=getcore.o
endif

F90=ftn
CC=cc

f.gnu: fhostone.F90 $(EXTRA)
    $(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.gnu
# $(F90) $(OPS) $(EXTRA) -fopenmp -fallow-argument-mismatch fhostone.F90 -O3 -o f.gnu
    rm -f getcore.o

c.gnu: phostone.c
    $(CC) $(OPS) -fopenmp phostone.c -O3 -o c.gnu

pp.gnu: ppong.c
    $(CC) $(OPS) $(WES) ppong.c -O3 -o pp.gnu

clean:
    rm -rf *o *mod* f.gnu c.gnu pp.gnu
```


makefile

PrgEnv-intel

```
SHELL:=/usr/bin/bash

recurse:
    module purge                ; \
    module load craype-x86-spr  ; \
    module load intel           ; \
    module load PrgEnv-intel    ; \
    $(MAKE) -f $(firstword $(MAKEFILE_LIST)) both

both: f.intel c.intel pp.intel

#defines USEFAST
include makefile.include

ifeq ($(USEFAST),yes)
OPS=-DUSEFAST
EXTRA=getcore.o
endif

F90=ftn
CC=cc

f.intel: fhostone.F90 $(EXTRA)
    $(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.intel
    rm -f getcore.o

c.intel: phostone.c
    $(CC) $(OPS) -fopenmp phostone.c -O3 -o c.intel

pp.intel: ppong.c
    $(CC) $(OPS) $(WES) ppong.c -O3 -o pp.intel

clean:
    rm -rf *o *mod* f.intel c.intel pp.intel
```


PrgEnv-* modules

PrgEnv-*	Loads for make	Loads of run (may not be required)
PrgEnv-cray	<code>module load craype-x86-spr</code> <code>module load PrgEnv-cray</code>	<code>module load craype-x86-spr</code> <code>module load PrgEnv-cray</code>
PrgEnv-gnu	<code>module load craype-x86-spr</code> <code>module load PrgEnv-gnu</code> <code>module load gcc</code>	<code>module load craype-x86-spr</code> <code>module load PrgEnv-gnu</code> <code>module load gcc</code>
PrgEnv-intel	<code>module load craype-x86-spr</code> <code>module load intel</code> <code>module load PrgEnv-intel</code>	<code>module load craype-x86-spr</code> <code>module load intel</code> <code>module load PrgEnv-intel</code>

Example Output PrgEnv-cray

MPI VERSION MPI VERSION : CRAY MPICH version 8.1.23.5 (ANL base 3.4a2)
MPI BUILD INFO : Tue Nov 29 12:42 2022 (git hash 210ae8b)

task	thread	node name	first task	# on node	core
total time	7.004				
0000	0000	X1005C2S3B1N1	0000	0000	0000
0000	0005	X1005C2S3B1N1	0000	0000	0005
0000	0007	X1005C2S3B1N1	0000	0000	0007
0000	0002	X1005C2S3B1N1	0000	0000	0002
0000	0001	X1005C2S3B1N1	0000	0000	0001
0000	0004	X1005C2S3B1N1	0000	0000	0004
0000	0006	X1005C2S3B1N1	0000	0000	0006
0000	0003	X1005C2S3B1N1	0000	0000	0003
0001	0007	X1005C3S1B0N0	0001	0000	0007
0001	0002	X1005C3S1B0N0	0001	0000	0002
0001	0004	X1005C3S1B0N0	0001	0000	0004
0001	0005	X1005C3S1B0N0	0001	0000	0005
0001	0003	X1005C3S1B0N0	0001	0000	0003
0001	0001	X1005C3S1B0N0	0001	0000	0001
0001	0006	X1005C3S1B0N0	0001	0000	0006
0001	0000	X1005C3S1B0N0	0001	0000	0000
mpi_init	0 1696732937.5528	1696732937.6665		0.1137	
mpi_init	1 1696732937.5150	1696732937.6673		0.1523	

Example Output PrgEnv-intel

```
MPI VERSION MPI VERSION      : CRAY MPICH version 8.1.23.5 (ANL base 3.4a2)
MPI BUILD INFO : Tue Nov 29 13:44 2022 (git hash 210ae8b)
```

task	thread	node name	first task	# on node	core
total time 7.004					
0000	0006	X1005C2S3B1N1	0000	0000	0006
0000	0005	X1005C2S3B1N1	0000	0000	0005
0000	0002	X1005C2S3B1N1	0000	0000	0002
0000	0000	X1005C2S3B1N1	0000	0000	0000
0000	0004	X1005C2S3B1N1	0000	0000	0004
0000	0003	X1005C2S3B1N1	0000	0000	0003
0000	0001	X1005C2S3B1N1	0000	0000	0001
0000	0007	X1005C2S3B1N1	0000	0000	0007
0001	0006	X1005C3S1B0N0	0001	0000	0006
0001	0005	X1005C3S1B0N0	0001	0000	0005
0001	0001	X1005C3S1B0N0	0001	0000	0001
0001	0004	X1005C3S1B0N0	0001	0000	0004
0001	0003	X1005C3S1B0N0	0001	0000	0003
0001	0000	X1005C3S1B0N0	0001	0000	0000
0001	0002	X1005C3S1B0N0	0001	0000	0002
0001	0007	X1005C3S1B0N0	0001	0000	0007
mpi_init	0	1696732969.0933	1696732969.2385	0.1452	
mpi_init	1	1696732969.0089	1696732969.2390	0.2301	

Example Output PrgEnv-gnu

```
MPI VERSION MPI VERSION      : CRAY MPICH version 8.1.23.5 (ANL base 3.4a2)
MPI BUILD INFO : Tue Nov 29 12:42 2022 (git hash 210ae8b)
```

task	thread	node name	first task	# on node	core
total time		7.004			
0000	0001	X1005C2S3B1N1	0000	0000	0001
0000	0003	X1005C2S3B1N1	0000	0000	0003
0000	0005	X1005C2S3B1N1	0000	0000	0005
0000	0004	X1005C2S3B1N1	0000	0000	0004
0000	0006	X1005C2S3B1N1	0000	0000	0006
0000	0000	X1005C2S3B1N1	0000	0000	0000
0000	0002	X1005C2S3B1N1	0000	0000	0002
0000	0007	X1005C2S3B1N1	0000	0000	0007
0001	0000	X1005C3S1B0N0	0001	0000	0000
0001	0002	X1005C3S1B0N0	0001	0000	0002
0001	0001	X1005C3S1B0N0	0001	0000	0001
0001	0006	X1005C3S1B0N0	0001	0000	0006
0001	0004	X1005C3S1B0N0	0001	0000	0004
0001	0005	X1005C3S1B0N0	0001	0000	0005
0001	0007	X1005C3S1B0N0	0001	0000	0007
0001	0003	X1005C3S1B0N0	0001	0000	0003
mpi_init	0	1696732984.4138	1696732984.5317	0.1179	
mpi_init	1	1696732984.3726	1696732984.5323	0.1597	

Interactions between PrgvEnv-* and Intel

- Many "intel" modules most of which you can ignore
- Loading PrgvEnv-intel loads a particular Intel compiler module
- Should not try to load PrgvEnv-intel and Intel MPI at the same time
- For running Intel MPI there are only two choices of module load gcc that work.

Module	Contains					Who
intel-classic-mixed/2023.2.0	icc	icx	ifort	ifx	NO_mpicc	NREL
intel-oneapi-mixed/2023.2.0	icc	icx	ifort	ifx	NO_mpicc	NREL
intel-oneapi/2023.2.0	icc	icx	ifort	ifx	NO_mpicc	NREL
intel-classic/2023.2.0	icc	icx	ifort	ifx	NO_mpicc	NREL
intel-oneapi-compilers/2023.2.0	icc	icx	ifort	ifx	NO_mpicc	NREL
intel-oneapi-mpi/2021.10.0-intel	NO_icc	NO_icx	NO_ifort	NO_ifx	mpicc	NREL
intel/2023.2.0	icc	icx	ifort	ifx	NO_mpicc	NREL

Module Specific Help for "intel-oneapi-compilers/2023.2.0"

Name : intel-oneapi-compilers
Version: 2023.2.0
Target : icelake

Intel oneAPI Compilers. Includes: icc, icpc, ifort, icx, icpx, ifx, and dpcpp. LICENSE INFORMATION: By downloading and using this software, you agree to the terms and conditions of the software license agreements at <https://intel.ly/393Cij0>.

Module Specific Help for "intel-oneapi-mpi/2021.10.0-intel"

Name : intel-oneapi-mpi
Version: 2021.10.0
Target : icelake

Intel MPI Library is a multifabric message-passing library that implements the open-source MPICH specification. Use the library to create, maintain, and test advanced, complex applications that perform better on high-performance computing (HPC) clusters based on Intel processors. LICENSE INFORMATION: By downloading and using this software, you agree to the terms and conditions of the software license agreements at <https://intel.ly/393Cij0>.

Module Specific Help for "intel-classic/2023.2.0" (and all of the others)

2023.2.0

/nopt/nrel/apps/compilers/08-23/spack/opt/spack/linux-rhel8-icelake/gcc-8.4.0/intel-oneapi-compilers-2023.2.0-cqpelkddr7kvjjmbqgs5ypz27m2bgqgt/compiler/2023.2.0

This modulefile defines the system paths and environment variables needed to use the Intel 'classic' compilers icc, ifort and icpc on Cray XE, XC, and XE systems. This modulefile may be loaded as a standalone or as part of PrgEnv-intel, when the user will call these products using cc, ftn, and CC. If loaded as part of PrgEnv-intel, it cannot be unloaded individually, but it can be swapped for another version.

Loading "cray" versions of gcc after loading Intel can "kick out" Intel MPI

Can give you an unexpected mix of Intel and Cray software.

Module loaded with module load intel-oneapi-mpi xxxxxxx module load gcc/*	Purge First	gcc (12.2.0)	gcc/10.1.0	gcc/10.3.0	gcc/11.2.0	gcc/12.1.0	gcc/12.2.0	gcc/13.1.0
intel-oneapi-compilers	No	CrayMPI	IntelMPI	CrayMPI	CrayMPI	CrayMPI	CrayMPI	IntelMPI
intel-oneapi-compilers	Yes	IntelMPI	IntelMPI	IntelMPI	IntelMPI	IntelMPI	IntelMPI	IntelMPI
intel-oneapi	No	CrayMPI	IntelMPI	CrayMPI	CrayMPI	CrayMPI	CrayMPI	IntelMPI
intel-oneapi	Yes	IntelMPI	IntelMPI	IntelMPI	IntelMPI	IntelMPI	IntelMPI	IntelMPI

icc available / MPI version

module load **cray-mpich-abi/8.1.23**

- module load cray-mpich-abi/8.1.23
 - Replaces Intel MPI with Cray MPI at runtime
 - Useful for cases where you have a binary but not the source
 - In theory, also works with programs built with MPICH
 - The command ldd can be used to see what version of MPI is being called


```
[tkaiser2@x1000c0s0b0n0 216704]$ldd c.impi | grep libmpi.so
```

```
libmpi.so.12 => /nopt/nrel/apps/mpi/08-23/spack/opt/spack/linux-rhel8-icelake/intel-2021.10.0/intel-oneapi-mpi-2021.10.0-7iolquprezbcmmeapg7rlcwfsd4r3rc7/mpi/2021.10.0/lib/release/libmpi.so.12 (0x00007f0c71d21000)
```

```
[tkaiser2@x1000c0s0b0n0 216704]$srun -n 2 ./c.impi -F
```

```
MPI VERSION Intel(R) MPI Library 2021.10 for Linux* OS
```

task	thread	node name	first task	# on node	core
0000	0000	x1000c0s0b0n0	0000	0000	0051
0000	0001	x1000c0s0b0n0	0000	0000	0039
0001	0000	x1000c0s0b0n0	0000	0001	0103
0001	0001	x1000c0s0b0n0	0000	0001	0091

```
[tkaiser2@x1000c0s0b0n0 216704]$module load cray-mpich-abi/8.1.23
```

Lmod is automatically replacing "cray-mpich/8.1.23" with "cray-mpich-abi/8.1.23".

```
[tkaiser2@x1000c0s0b0n0 216704]$ldd c.impi | grep libmpi.so
```

```
libmpi.so.12 => /opt/cray/pe/mpich/8.1.23/ofc/crayclang/10.0/lib-abi-mpich/libmpi.so.12 (0x00007fae15df4000)
```

```
[tkaiser2@x1000c0s0b0n0 216704]$srun -n 2 ./c.impi -F
```

```
MPI VERSION MPI VERSION : CRAY MPICH version 8.1.23.5 (ANL base 3.4a2)
```

```
MPI BUILD INFO : Tue Nov 29 12:42 2022 (git hash 210ae8b)
```

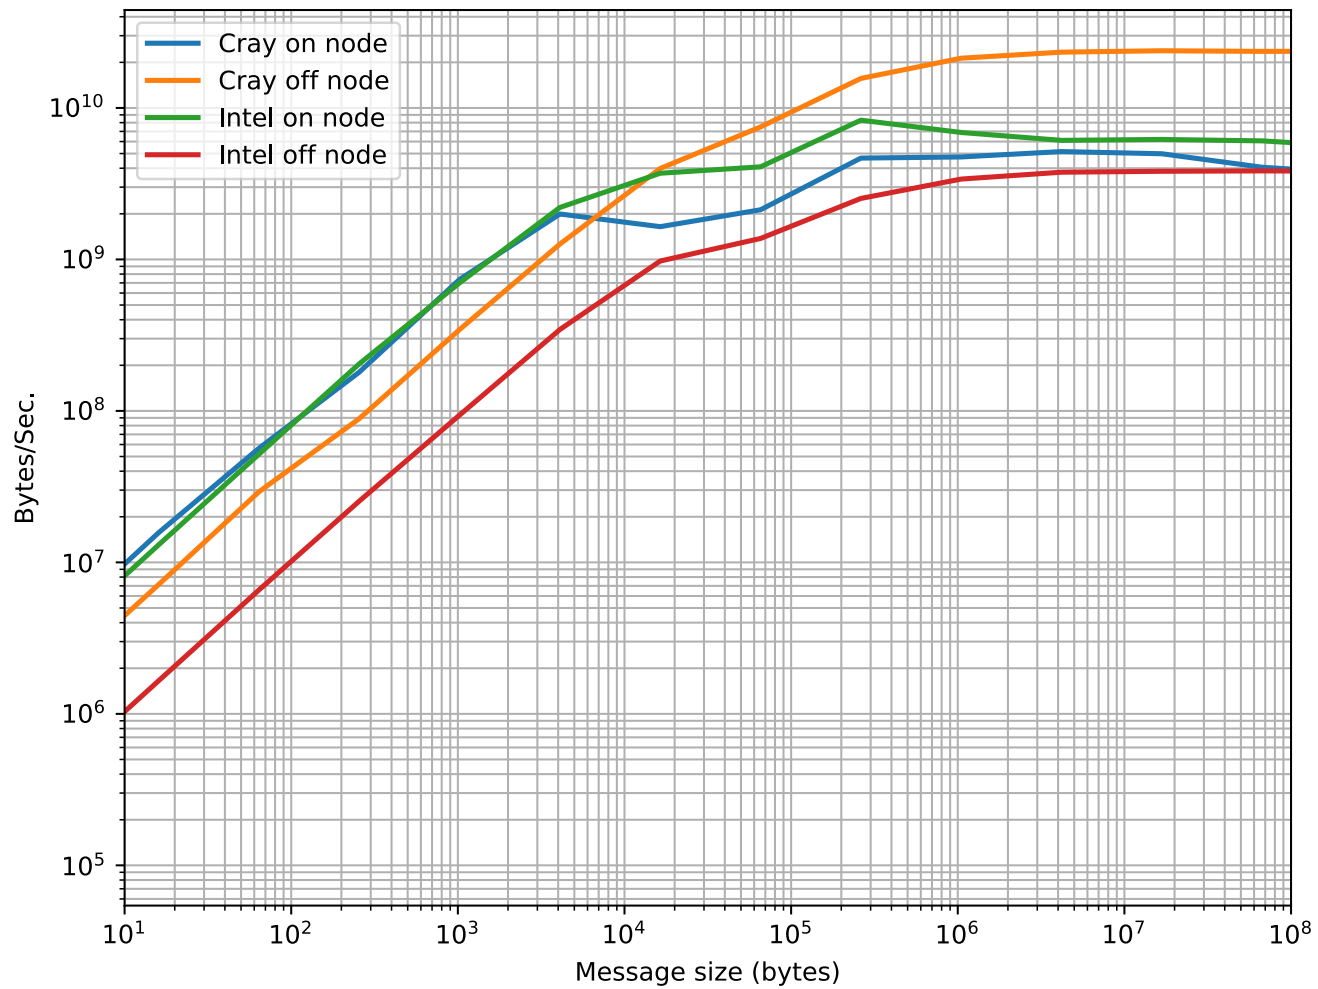
task	thread	node name	first task	# on node	core
0000	0000	x1000c0s0b0n0	0000	0000	0051
0000	0001	x1000c0s0b0n0	0000	0000	0039
0001	0000	x1000c0s0b0n0	0000	0001	0103
0001	0001	x1000c0s0b0n0	0000	0001	0093

```
[tkaiser2@x1000c0s0b0n0 216704]$
```


BACK TO PINGPONG

- Measures message speed as a function of size
- Measure on node and off node
- All PrgEnv-* will run at the same speed
- We compare Intel-MPI to PrgEnv-*
- Which is better:
 - Intel in general better on node
 - PrgEnv-* is better off node

Kestrel MPI Bandwidth



AFFINITY & WHY IMPORTANT


- Affinity - mapping of threads/tasks to cores

- Kestrel 104 cores/node

- 2 chips (Intel 8470)

- 52 cores each

- 4 "tiles" with 13 cores each



- Worst case: Multiple threads/tasks can end up on the same core potentially reducing performance by 2X or maybe much more
- Also: You may want to put threads/tasks on particular tiles to maximize communications or memory access
- Possible to have different MPI tasks to have different # threads (Ask if interested)

CUT TO THE CHASE:

- With the proper setting in sbatch scripts and the srun command we are able to “trivially” get apps to behave reasonably for all cases I tested on TDS, Swift and Eagle.
- Masks (a mapping list) allow a fine grain placement of tasks & threads to cores

SUFFICIENT EXAMPLE

For 2 nodes, 18 tasks per node, 2 threads per task

```
#!/usr/bin/bash
```

```
#SBATCH --nodes=2
```

```
#SBATCH --exclusive
```

```
#SBATCH --export=ALL
```

```
export OMP_PLACES=cores
```

```
export OMP_PROC_BIND=spread
```

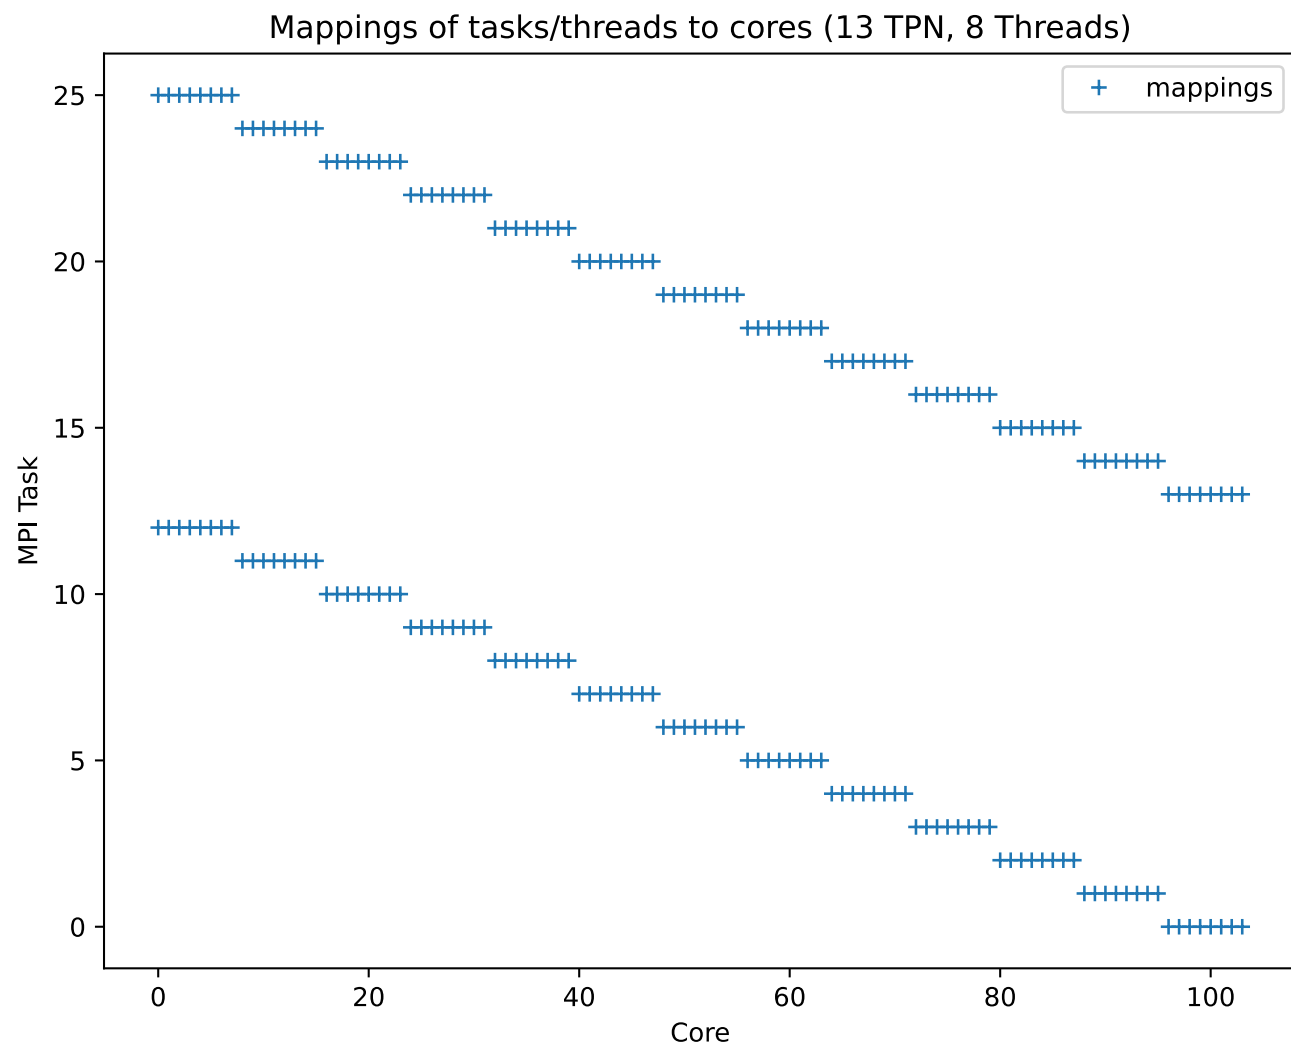
```
export OMP_NUM_THREADS=13
```

```
srun --mpi=pmi2 --threads-per-core=1 --tasks-per-node=4 --cpus-per-task=13 ...
```


With the setting show in red we get good mappings for Kestrel, Eagle and Swift


`--cpus-per-task = OMP_NUM_THREADS`

Add `-cpu-bind=v` to see the binding


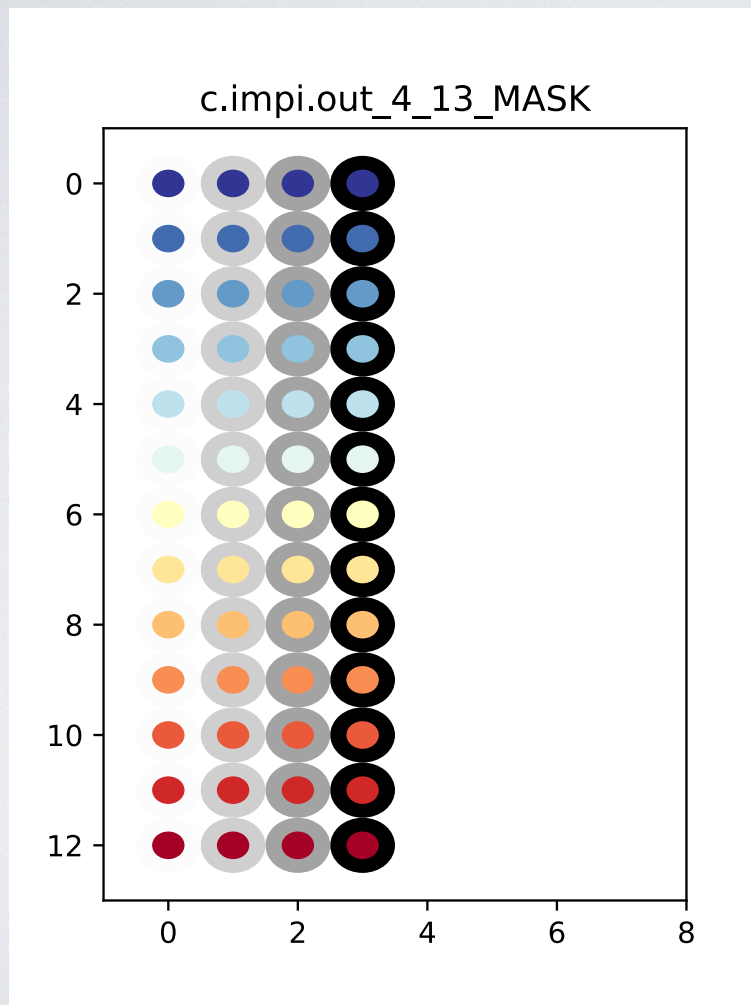


Mapping of tasks/threads to cores (13 TPN, 8 Threads)





EXAMPLE 4 & 8 TASKS/NODE 13 THREADS



Tasks/node	Threads/task	Cores/node
104	1	104
52	2	104
52	1	52
26	4	104
26	2	52
26	1	26
13	4	52
13	2	26
13	1	13
8	13	104
4	26	104
4	13	52
4	8	32
2	52	104
2	26	52
2	13	26
2	8	16
1	104	104
1	91	91
1	78	78
1	65	65
1	52	52
1	26	26
1	13	13
1	8	8

Tests run

MPI/Compiler sets

- PrgEnv-cray
- PrgEnv-gnu
- PrgEnv-intel
- intel-oneapi

OUR EXAMPLE SCRIPT

- Originally designed as an compile/run and affinity tester
- Sets up a new directory and goes there
- Copies all required files
- Does a make for all versions
- Loops over # MPI tasks and # OpenMP threads (input file cases)
 - Loops over two types of thread binding (scattered and manual)
 - Steps through MPI versions with both C and Fortran

OUR EXAMPLE SCRIPT

- Reports lots of information
 - Bindings for each run
 - Normal program output for phostrun includes mapping of tasks and threads to nodes and cores
 - MPI launch times
- Runs a single instance of ppong for each version of MPI (runs the setup script todo.py to create an input file)
- Final output is a report of successful/failed mapping
 - Success = expected unique combinations of nodes and cores
 - Good News: It works for all tested versions of MPI and mappings

OVERKILL FOR MOST PEOPLE

- While you can run this for the full set you might not want to use the allocation hours (minutes)
- Suggested use...
 - Run phostone using the exact run arguments you use for your production code to see how it maps tasks and threads to cores


```
#!/usr/bin/bash
#SBATCH --job-name="affinity"
#SBATCH --nodes=2
#SBATCH --exclusive
#SBATCH --export=ALL
#SBATCH --time=04:00:00
#SBATCH --partition=standard

BASE=`pwd`

#Make a new directory and go there
STDIR=`pwd`
mkdir $SLURM_JOB_ID
cd $SLURM_JOB_ID

#optionally wait between launches
mywait () { sleep 0; }

#Copy everything
printenv > env
cat $0 > script

cp $BASE/make* .
cp $BASE/Makefile .
cp $BASE/fhostone.F90 .
cp $BASE/phostone.c .
cp $BASE/cases .
cp $BASE/post .
cp $BASE/ppong.c .
cp $BASE/getcore.c .
cp $BASE/maskgenerator.py .
cp $BASE/todo.py .
cp $BASE/tymer .

tar -czf recreate.tgz *

#Create input for ppong
./todo.py

#Build our programs
make all > make.log 2>&1
make pp > make.pp 2>&1

#Command line arguments for phostone
CLA="-i -F -E -t 7"
export FEFE=f
export CEFE=c
```

Setup and "make"

make*	Makefiles
Makefile	Driver Makefile
cases	File of tasks and threads
post	Post processing script
maskgenerator.py	Manually creates mapping of threads to cores
todo.py	Creates input file for ppong
tymer	Nice wall clock timer

Makefile (full)

```
all : impi  cray  gnu  intel  open  mpich  openg  mpichg  dmod
```

```
impi: makeimpi
     make -f makeimpi
```

```
cray: makeprgcray
     make -f makeprgcray
```

```
gnu: makeprggnu
     make -f makeprggnu
```

```
intel: makeprgintel
       make -f makeprgintel
```

```
open: makeopen
       make -f makeopen
```

```
mpich: makempich
       make -f makempich
```

```
openg: makeopen_g
       make -f makeopen_g
```

```
mpichg: makempich_g
        make -f makempich_g
```

```
clean:
    make -f makeimpi clean
    make -f makeprgintel clean
    make -f makeprggnu clean
    make -f makeprgcray clean
    make -f makeopen clean
    make -f makempich clean
    make -f makeopen_g clean
    make -f makempich_g clean
    rm -rf runall.tgz simple.tgz
```

```
dmod:
     rm -rf *.o *mod
```

```
pp: pp.impi pp.cray pp.gnu pp.intel pp.open pp.oneapi pp.mpich pp.openg pp.mpichg
```

```
pp.impi: makeimpi
         make -f makeimpi pp.impi
```

```
pp.cray: makeprgcray
         make -f makeprgcray pp.cray
```

```
pp.gnu: makeprggnu
        make -f makeprggnu pp.gnu
```

```
pp.intel: makeprgintel
          make -f makeprgintel pp.intel
```

```
pp.open: makeopen
          make -f makeopen pp.open
```

```
pp.mpich: makempich
          make -f makempich pp.mpich
```

```
pp.openg: makeopen_g
          make -f makeopen_g pp.openg
```

```
pp.mpichg: makempich_g
           make -f makempich_g pp.mpichg
```

```
tar:
    tar -czf runall.tgz \
        cases eagle ecases fhostone.F90 getcore.c makelapi Makefile makefile.include \
        makeimpi makeopen makeprgcray makeprggnu makeprgintel maskgenerator.py masks.txt \
        phostone.c post ppong.c readme.md runall runpp subsweep sweep todo.py tymer \
        scases array mapping.py simple makempich makempich_g makeopen_g
```

```
simple.tgz:
    tar -czf simple.tgz fhostone.F90 getcore.c makelapi Makefile makefile.include \
        makefile.org makeimpi makeopen makeprgcray makeprggnu makeprgintel \
        phostone.c post ppong.c simple makempich makempich_g makeopen_g
```



```

#LOOPING
export CRAY_OMP_CHECK_AFFINITY=TRUE
export nc=`cat cases | wc -l`
for il in `seq $nc` ; do
    aline=`cat cases | head -$il | tail -1`
    ntpn=`echo $aline | awk {'print $1'}`
    nthrd=`echo $aline | awk {'print $2'}`
    export OMP_NUM_THREADS=$nthrd
    for bindit in NONE MASK ; do
        #export KMP_AFFINITY=scatter
        export OMP_PROC_BIND=spread
        export BIND=--cpu-bind=v,${bindit}
        unset CPUS_TASK
        if [ $bindit == MASK ] ; then
            cores=`expr $ntpn \* $nthrd`
            MASK=`./maskgenerator.py $cores $ntpn`
            BIND="--cpu-bind=v,mask_cpu:$MASK"
        fi
        if [ $bindit == NONE ] ; then
            BIND="--cpu-bind=v"
        fi
        export CPUS_TASK="--cpus-per-task=$nthrd"
        echo $ntpn $nthrd >> srunsettings
        echo $BIND $CPUS_TASK >> srunsettings
        printenv | egrep "OMP_|KMP_" >> srunsettings
        echo --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK >> srunsettings
    done
done

```

Looping

- Get a task/thread count from each line of cases
- We try two types of thread binding, spread (NONE) and manually (MASK)
 - The script maskgenerator.py creates a string describing a mapping of tasks/threads to cores
 - This is passed to run using the --cpu-bind option
 - Save information for each iteration


```

./tymer mytimes PrgEnv-intel
    module purge
    module load craype-x86-spr
    module load intel
    module load PrgEnv-intel

./tymer mytimes fortran
    mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./FEXE.intel $CLA > f.intel.out_${ntpn}_${nthrd}_${bindit} \
        2> f.intel.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes c
    mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./CEXE.intel $CLA > c.intel.out_${ntpn}_${nthrd}_${bindit} \
        2> c.intel.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished

    if [[ $nthrd -eq 1 && $ntpn -eq 104 && $bindit == NONE ]] ; then
    mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./pp.intel $CLA > pp.intel.xxx_${ntpn}_${nthrd}_${bindit} \
        2> pp.intel.iii_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished ppong
    fi

./tymer mytimes PrgEnv-gnu
    module purge
    module load craype-x86-spr
    module load PrgEnv-gnu

./tymer mytimes fortran
    mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./FEXE.gnu $CLA > f.gnu.out_${ntpn}_${nthrd}_${bindit} \
        2> f.gnu.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes c
    mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./CEXE.gnu $CLA > c.gnu.out_${ntpn}_${nthrd}_${bindit} \
        2> c.gnu.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished

    if [[ $nthrd -eq 1 && $ntpn -eq 104 && $bindit == NONE ]] ; then
    mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./pp.gnu $CLA > pp.gnu.xxx_${ntpn}_${nthrd}_${bindit} \
        2> pp.gnu.iii_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished ppong
    fi

```



```

./tymer mytimes PrgEnv-cray
  module purge
  module load craype-x86-spr
  module load PrgEnv-cray

./tymer mytimes fortran
  mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./FEXE.cray $CLA > f.cray.out_${ntpn}_${nthrd}_${bindit} \
    2> f.cray.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes c
  mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./CEXE.cray $CLA > c.cray.out_${ntpn}_${nthrd}_${bindit} \
    2> c.cray.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished

  if [[ $nthrd -eq 1 && $ntpn -eq 104 && $bindit == NONE ]] ; then
  mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./pp.cray $CLA > pp.cray.xxx_${ntpn}_${nthrd}_${bindit} \
    2> pp.cray.iii_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished ppong
  fi

./tymer mytimes intel-oneapi
  module purge
  module load intel-oneapi
  module load libfabric

./tymer mytimes fortran
  mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./FEXE.impi $CLA > f.impi.out_${ntpn}_${nthrd}_${bindit} \
    2> f.impi.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes c
  mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./CEXE.impi $CLA > c.impi.out_${ntpn}_${nthrd}_${bindit} \
    2> c.impi.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished

  if [[ $nthrd -eq 1 && $ntpn -eq 104 && $bindit == NONE ]] ; then
  mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./pp.impi $CLA > pp.impi.xxx_${ntpn}_${nthrd}_${bindit} \
    2> pp.impi.iii_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished ppong
  fi
done
done

```



```
./post  
. ./post | sort -n > posit  
getstate postit nope > report  
getstate postit worked >> report  
mv $STDIR/slurm-$SLURM_JOB_ID.out .
```

Post processing

- Report of successful and failed phostone runs
- Copy slurm output to our final directory
- Might want to look at output from ppong
 - Bandwidth
 - MPI_Barrier rate
- Might want to look at MPI_Init times from phostone

REPORT

```
tkaiser2-37907s:177402 tkaiser2$ cat report
```

```
c.cray.out_104_1_MASK 208 208
c.cray.out_104_1_NONE 208 208
c.cray.out_1_104_MASK 208 208
c.cray.out_1_104_NONE 208 208
c.cray.out_1_8_MASK 16 16
c.cray.out_1_8_NONE 16 16
c.gnu.out_104_1_MASK 208 208
c.gnu.out_104_1_NONE 208 208
c.gnu.out_1_104_MASK 208 208
c.gnu.out_1_104_NONE 208 208
c.gnu.out_1_8_MASK 16 16
c.gnu.out_1_8_NONE 16 16
c.impi.out_104_1_MASK 208 208
c.impi.out_104_1_NONE 208 208
c.impi.out_1_104_MASK 208 208
c.impi.out_1_104_NONE 208 208
c.impi.out_1_8_MASK 16 16
c.impi.out_1_8_NONE 16 16
c.intel.out_104_1_MASK 208 208
c.intel.out_104_1_NONE 208 208
c.intel.out_1_104_MASK 208 208
c.intel.out_1_104_NONE 208 208
c.intel.out_1_8_MASK 16 16
c.intel.out_1_8_NONE 16 16
```

```
f.cray.out_104_1_MASK 208 208
f.cray.out_104_1_NONE 208 208
f.cray.out_1_104_MASK 208 208
f.cray.out_1_104_NONE 208 208
f.cray.out_1_8_MASK 16 16
f.cray.out_1_8_NONE 16 16
f.gnu.out_104_1_MASK 208 208
f.gnu.out_104_1_NONE 208 208
f.gnu.out_1_104_MASK 208 208
f.gnu.out_1_104_NONE 208 208
f.gnu.out_1_8_MASK 16 16
f.gnu.out_1_8_NONE 16 16
f.impi.out_104_1_MASK 208 208
f.impi.out_104_1_NONE 208 208
f.impi.out_1_104_MASK 208 208
f.impi.out_1_104_NONE 208 208
f.impi.out_1_8_MASK 16 16
f.impi.out_1_8_NONE 16 16
f.intel.out_104_1_MASK 208 208
f.intel.out_104_1_NONE 208 208
f.intel.out_1_104_MASK 208 208
f.intel.out_1_104_NONE 208 208
f.intel.out_1_8_MASK 16 16
f.intel.out_1_8_NONE 16 16
```


WE ARE SKIPPING SOME INSTALLED VERSIONS OF MPI

```
openmpi/4.1.5-gcc  
openmpi/4.1.5-intel  
mpich/4.1-gcc  
mpich/4.1-intel
```

- Skipped
- Don't perform as well
- Running a single instance of these usually works.
- Running many in succession has a bad habit of hanging


```
SHELL:=/usr/bin/bash
```

```
recurse:
```

```
    module purge                ; \
    module load mpich/4.1-gcc    ; \
    module load gcc/13.1.0       ; \
    $(MAKE) -f $(firstword $(MAKEFILE_LIST)) both
```

```
# To use intelversion of the compilers
# replace thw two lines above with these
# module load mpich/4.1-intel ; \
# module load intel-oneapi ; \
# You should replace the same lines in the run script.
```

```
both: f.mpichg c.mpichg pp.mpichg
```

```
#defines USEFAST
include makefile.include
```

```
ifeq ($(USEFAST),yes)
#OPS=-DUSEFAST
#EXTRA=getcore.o
endif
```

```
F90=mpif90
CC=mpicc -lm
```

```
f.mpichg: fhostone.F90 $(EXTRA)
    $(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.mpichg
    rm -f getcore.o
```

```
c.mpichg: phostone.c
    $(CC) $(OPS) -fopenmp phostone.c -O3 -o c.mpichg
```

```
pp.mpichg: ppong.c
    $(CC) $(OPS) $(WES) ppong.c -O3 -o pp.mpichg
```

```
clean:
    rm -rf *o *mod* f.mpichg c.mpichg pp.mpichg
```

mpich/4.1-gcc

mpich/4.1-intel

```
SHELL:=/usr/bin/bash

recurse:
    module purge                ; \
    module load mpich/4.1-intel ; \
    module load intel-oneapi    ; \
    $(MAKE) -f $(firstword $(MAKEFILE_LIST)) both

# To use gcc version of the compilers
# replace the two lines above with these
# module load mpich/4.1-gcc ; \
# module load gcc ;\
# You should replace the same lines in the run script.

both: f.mpich c.mpich pp.mpich

#defines USEFAST
include makefile.include

ifeq ($(USEFAST),yes)
#OPS=-DUSEFAST
#EXTRA=getcore.o
endif

F90=mpif90
CC=mpicc -lm

f.mpich: fhostone.F90 $(EXTRA)
    $(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.mpich
    rm -f getcore.o

c.mpich: phostone.c
    $(CC) $(OPS) -fopenmp phostone.c -O3 -o c.mpich

pp.mpich: ppong.c
    $(CC) $(OPS) $(WES) ppong.c -O3 -o pp.mpich

clean:
    rm -rf *o *mod* f.mpich c.mpich pp.mpich
```



```
SHELL:=/usr/bin/bash
```

```
recurse:
    module purge                ; \
    module load openmpi/4.1.5-intel ; \
    module load intel-oneapi    ; \
    $(MAKE) -f $(firstword $(MAKEFILE_LIST)) both
```

```
# To run with gcc /Openmpi replace the lines above with these
# module load openmpi/4.1.5-gcc ; \
# module load gcc      ; \
# You should also replace the lines in the run script.
```

```
both: f.open c.open pp.open
```

```
#defines USEFAST
include makefile.include
```

```
ifeq ($(USEFAST),yes)
#OPS=-DUSEFAST
#EXTRA=getcore.o
endif
```

```
F90=mpif90
CC=mpicc -lm
```

```
f.open: fhostone.F90 $(EXTRA)
    $(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.open
    rm -f getcore.o
```

```
c.open: phostone.c
    $(CC) $(OPS) -fopenmp phostone.c -O3 -o c.open
```

```
pp.open: ppong.c
    $(CC) $(OPS) $(WES) ppong.c -O3 -o pp.open
```

```
clean:
    rm -rf *o *mod* f.open c.open pp.open
```

openmpi/4.1.5-intel


```
SHELL:=/usr/bin/bash
```

```
recurse:
```

```
module purge                ; \  
module load openmpi/4.1.5-gcc ; \  
module load gcc/13.1.0      ; \  
$(MAKE) -f $(firstword $(MAKEFILE_LIST)) both
```

```
# To run with intel / Openmpi replace the lines above with these  
# module load openmpi/4.1.5-intel ; \  
# module load intel-oneapi ; \  
# You should also replace the lines in the run script.
```

```
both: f.openg c.openg pp.openg
```

```
#defines USEFAST  
include makefile.include
```

```
ifeq ($(USEFAST),yes)  
#OPS=-DUSEFAST  
#EXTRA=getcore.o  
endif
```

```
F90=mpif90  
CC=mpicc -lm
```

```
f.openg: fhostone.F90 $(EXTRA)  
$(F90) $(OPS) $(EXTRA) -fopenmp fhostone.F90 -O3 -o f.openg  
rm -f getcore.o
```

```
c.openg: phostone.c  
$(CC) $(OPS) -fopenmp phostone.c -O3 -o c.openg
```

```
pp.openg: ppong.c  
$(CC) $(OPS) $(WES) ppong.c -O3 -o pp.openg
```

```
clean:  
rm -rf *o *mod* f.openg c.openg pp.openg
```

openmpi/4.1.5-gcc

:<<SKIP

```
tymer mytimes openmpi/4.1.5-gcc
module purge
module load openmpi/4.1.5-gcc
module load gcc
mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./FEXE.openg $CLA > f.openg.out_${ntpn}_${nthrd}_${bindit} \
    2> f.openg.info_${ntpn}_${nthrd}_${bindit}
mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./CEXE.openg $CLA > c.openg.out_${ntpn}_${nthrd}_${bindit} \
    2> c.openg.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished

if [[ $nthrd -eq 1 && $ntpn -eq 104 && $bindit == NONE ]] ; then
mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./ppong.openg $CLA > pp.openg.xxx_${ntpn}_${nthrd}_${bindit}
    2> pp.openg.iii_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished ppong
f1

tymer mytimes openmpi/4.1.5-intel
module purge
module load openmpi/4.1.5-intel
module load intel-oneapi
mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./FEXE.open $CLA > f.open.out_${ntpn}_${nthrd}_${bindit} \
    2> f.open.info_${ntpn}_${nthrd}_${bindit}
mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./CEXE.open $CLA > c.open.out_${ntpn}_${nthrd}_${bindit} \
    2> c.open.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished

if [[ $nthrd -eq 1 && $ntpn -eq 104 && $bindit == NONE ]] ; then
mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./ppong.open $CLA > pp.open.xxx_${ntpn}_${nthrd}_${bindit}
    2> pp.iopen.iii_${ntpn}_${nthrd}_${bindit}
```



```

./tymer mytimes finished ppong
fi
tymer mytimes mpich/4.1-intel
module purge
module load mpich/4.1-intel
module load intel-oneapi
module load libfabric
unset UCX_NET_DEVICES
mywait; srun --mpi=pmi2 $BIND --time=00:03:00 --tasks-per-node=$ntpn $CPUS_TASK ./FEXE.mpich $CLA > f.mpich.out_${ntpn}_${nthrd}_${bindit} \
2> f.mpich.info_${ntpn}_${nthrd}_${bindit}
mywait; srun --mpi=pmi2 $BIND --time=00:03:00 --tasks-per-node=$ntpn $CPUS_TASK ./CEXE.mpich $CLA > c.mpich.out_${ntpn}_${nthrd}_${bindit} \
2> c.mpich.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished

if [[ $nthrd -eq 1 && $ntpn -eq 104 && $bindit == NONE ]] ; then
mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./ppong.mpich $CLA > pp.mpich.xxx_${ntpn}_${nthrd}_${bindit}
2> pp.mpich.iii_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished ppong
fi

tymer mytimes mpich/4.1-gcc
module purge
module load mpich/4.1-gcc
module load gcc
module load libfabric
unset UCX_NET_DEVICES
mywait; srun --mpi=pmi2 $BIND --time=00:03:00 --tasks-per-node=$ntpn $CPUS_TASK ./FEXE.mpichg $CLA > f.mpichg.out_${ntpn}_${nthrd}_${bindit} \
2> f.mpichg.info_${ntpn}_${nthrd}_${bindit}
mywait; srun --mpi=pmi2 $BIND --time=00:03:00 --tasks-per-node=$ntpn $CPUS_TASK ./CEXE.mpichg $CLA > c.mpichg.out_${ntpn}_${nthrd}_${bindit} \
2> c.mpichg.info_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished

if [[ $nthrd -eq 1 && $ntpn -eq 104 && $bindit == NONE ]] ; then
mywait; srun --mpi=pmi2 $BIND --tasks-per-node=$ntpn $CPUS_TASK ./ppong.mpichg $CLA > pp.mpichg.xxx_${ntpn}_${nthrd}_${bindit}
2> pp.mpichg.iii_${ntpn}_${nthrd}_${bindit}
./tymer mytimes finished ppong
fi

```

SKIP